
15 seconds.

That's how long my AI agent takes to read, analyse & reply to a support email.

Fully autonomous. Zero human needed.

I built this. Here's how it works. →

31

Tickets

<15s

Per Email

70

Tests Passed

The Problem

Support teams are drowning.

- Same questions asked 100x a day
 - Avg email response time: 12+ hours
 - Customers write in Hindi, German, Japanese...
 - Agents burn out on repetitive work
 - No visibility — what issues are trending?
-

There had to be a better way.

The Solution

AI agent that handles emails end-to-end.

- Auto-detects language — replies in customer's language
- RAG searches product manual for answers
- Max 2 questions — polite limit message for 3+
- Auto-generates subject for blank emails
- Blocks email loops — 3-layer protection
- Live dashboard — 13 charts, 7 KPIs, 6 filters

31

Processed

65%

Resolved

100%

SLA

Architecture

5-Phase AI Pipeline — All Complete

- 01 **Email Trigger** — Gmail → N8N workflow fires
- 02 **AutoGen AI** — 10-step analysis pipeline
- 03 **RAG + GPT-4o** — Pinecone top_k=8 + GPT-4o reply
- 04 **Smart Routing** — Auto-reply OR escalate to human
- 05 **Live Dashboard** — 13 charts · 7 KPIs · 6 filters · Search · CSV

N8N · AutoGen · GPT-4o · Pinecone · Flask · Chart.js · Sheets · Telegram

AutoGen AI Engine

10 steps. Every single email.

- 1 **Self-Reply Block** — Prevents email loops at source
- 2 **Auto-Subject** — GPT-4o generates if blank
- 3 **FAQ Cache** — 24ms for repeated questions
- 4 **Spam/OOO** — 25+ keyword patterns
- 5 **Language Detect** — ISO code: hi/es/de/ja/ar...
- 6 **Question Limit** — Max 2 per email, polite message
- 7 **RAG Query** — top_k=8 cosine similarity
- 8 **GPT-4o Reply** — temp=0.3, max_tokens=2000
- 9 **Escalation** — 3 triggers
- 10 **Cache Save** — confidence>=0.5

Live Dashboard v2.0

72.61.178.139:8080 · Auto-refresh 30s

31

Tickets

65%

Resolution

0.33

RAG Score

100%

SLA

- **13 Charts** — Fishbone·Donuts·Bar·Line·Gauge + values+%
- **Search** — Updates ALL charts — not just table
- **6 Filters** — Priority·Sentiment·Status·Language·Today/Custom
- **Export CSV** — Download filtered tickets one click
- **Dynamic Language** — Auto-adds HI/ES/DE/JA to filter
- **RAG Trend** — Live from rag_score_at_escalation

10+ Languages

Auto-detected. Auto-replied.

Hindi	HI
Spanish	ES
German	DE
Japanese	JA
Arabic	AR
French	FR
Mixed Lang	DOMINANT

Replies in dominant language for mixed emails

Real Results

Numbers don't lie.

200x

Faster on Cache

<15s

End-to-End

70

Tests Passed

- 10+ languages auto-detected and replied correctly
- 3 escalation triggers — never misses a case
- 6 Google Sheets tabs auto-populated in real-time
- Email loop prevented with 3-layer protection
- Customer follow-up replies handled automatically

Smart Escalation

It knows when NOT to reply.

T1 **Confidence<0.4** — RAG cannot find answer

T2 **Angry+High Priority** — Sentiment + priority combo

T3 **Keywords** — refund/lawyer/broken/kaputt

ESCALATION — Human Review Needed!

Priority: Very High | Confidence: Low (0.327) | Sentiment: Angry

Subject: Demand refund — lawyer involved

Case Reference: 1780121056189

Want to build something like this?

- N8N — workflow automation
- AutoGen + FastAPI — 10-step AI engine
- OpenAI GPT-4o + Embeddings
- Pinecone — vector database (top_k=8)
- Flask + Chart.js — 13-chart live dashboard
- Google Sheets + Gmail + Telegram

70 test cases passed · 10+ languages · <15s per email

Drop a comment — Follow for more AI agent builds